



Convex Optimization

Project 1



Spring 1401

Due date: 30th of Ordibehesht

Predicting complete rankings. A (complete) ranking of K items consists of an ordering of the items from rank 1 to rank K . For example, these could be K candidates, ranked from 1 (best) to K (worst), or the order in which K horses cross the finish line in a race. We represent a ranking of K items as a vector $\pi \in \mathbf{R}^K$, with π_i the rank of item i . In the vector π , the numbers $1, \dots, K$ each appear exactly once (so it can also be considered a permutation), so there are $K!$ different rankings. We will let $\mathcal{P} \subset \mathbf{R}^K$ denote the set of all $K!$ rankings.

For example with $K = 3$, $(2, 3, 1)$ and $(1, 3, 2)$ are two of the six possible rankings. In the first ranking, item 1 has rank 2, whereas in the second ranking, item 1 has rank 1. Both rankings agree that item 2 has rank 3.

There are many ways to assign a distance between two rankings π and σ , but we will use a simple one, $(1/2)\|\pi - \sigma\|_1$. This distance is zero if and only if $\pi = \sigma$, and one if and only if π and σ assign the same rank to all items except two, whose ranks are off by one. The maximum possible distance is $K^2/4$ for K even and $(K^2 - 1)/4$ for K odd, achieved by, e.g., $\pi = (1, 2, \dots, K)$ and $\sigma = (K, K-1, \dots, 1)$. The average distance between two randomly chosen rankings is $(K^2 - 1)/6$. (These observations are not relevant for this problem, but only meant to give you an idea of the range and scale of the distance between rankings.)

We wish to build a predictor of an outcome which is a ranking, based on a vector of features. We denote the predictor as $P : \mathbf{R}^d \rightarrow \mathcal{P}$, where $P(x)$ is the ranking we predict when the feature vector is $x \in \mathbf{R}^d$. We will judge a predictor by the average distance between the true ranking and the predicted one, on a test set of data $(x_i^{\text{test}}, \pi_i^{\text{test}})$, $i = 1, \dots, N^{\text{test}}$ (that presumably was not used to develop or fit the predictor):

$$\frac{1}{2N^{\text{test}}} \sum_{i=1}^{N^{\text{test}}} \|\pi_i^{\text{test}} - P(x_i^{\text{test}})\|_1.$$

We refer to this quantity as the average test error of the predictor. (The smaller this is, the better the predictor performs on the test data set.) We will consider a simple predictor of the form $P(x) = \Pi(\theta x)$, where $\theta \in \mathbf{R}^{K \times d}$ is the predictor coefficient matrix, and $\Pi : \mathbf{R}^K \rightarrow \mathcal{P}$ is Euclidean projection onto \mathcal{P} . (We will describe this projection in more detail below, but for now we note that if there are multiple rankings that are closest to θx , we arbitrarily choose one.)

We choose the predictor parameter matrix θ to minimize

$$\frac{1}{2N} \sum_{i=1}^N \|\pi_i - \theta x_i\|_1,$$

where (x_i, π_i) , $i = 1, \dots, N$, is some given training data. (Note that this objective would become the average distance between the true and predicted rankings if we replace θx_i with $\Pi(\theta x_i)$, but then the objective is no longer convex.)

Projection onto rankings. You can use the following, without deriving or justifying it. The projection $\pi = \Pi(y)$ is the vector of rank orders of the entries of y in nondecreasing order. For example with $y = (1.1, -0.3, 0.5, 0.4)$, we have $\Pi(y) = (4, 1, 3, 2)$, since the first entry of y is the largest (i.e., has rank 4), the second entry of y is the smallest (i.e., has rank 1), and so on. So we can compute $\Pi(y)$ by sorting the entries of y (breaking any ties arbitrarily), keeping track of the sort ordering.

Explain how to fit the predictor using the training data with convex optimization.

The data file `ranking_est_data.*` contains functions that generate synthetic training and test data, as well as a function that implements Π . The data are in the matrices X_{train} , π_{train} , X_{test} , π_{test} , and the projection Π is given in `Pi()`. Fit the predictor using the training data, and give the average distance between the true and predicted ranking on both the training and test data sets.

Good Luck!